CUDA 프로그래밍

CUDA Programming

biztripcru@gmail.com

© 2021-2022. biztripcru@gmail.com. All rights reserved. 모든 저작권은 biztripcru@gmail.com 에게 있습니다.



Matrix Copy

행렬 복사



본 동영상과, 본 동영상 촬영에 사용된 발표 자료는 저작권법의 보호를 받습니다. 본 동영상과 발표 자료는 공개/공유/복제/상업적 이용 등, **개인 수강 이외의 다른 목적으로 사용하지 못합니다.**

© 2021-2022. biztripcru@gmail.com. All rights reserved. 모든 저작권은 biztripcru@gmail.com 에게 있습니다.

cuda-chap-22-mat-copy_220913.pptx

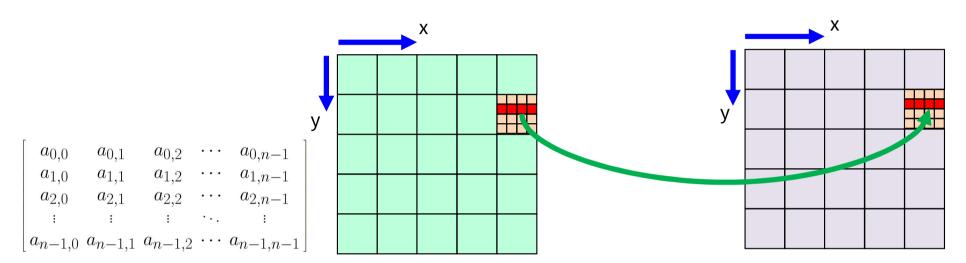
내용contents

- matrix copy theoretical limit
 - CPU version
 - memcpy version
 - CUDA naïve version global memory
 - CUDA shared memory version
 - CUDA memcpy2D
 - naive: (전문적) 지식이 없는

Matrix Copy

- simply copy a matrix to another
 - C[i,j] = A[i,j]
 - pitched matrices for the best performance
 - for simplicity, we assume square matrices
- 다른 matrix 연산의 입장에서는?
 - theoretical limit 이론적 한계 for matrix operations
 - the best score ^{최고 기록} for any matrix operations
 - 이 기록에 접근할 수록 최적화가 잘 되었음
 - 이 기록을 추월했다면, 뭔가 잘못 되었음...

Matrix Copy - CPU version



2D matrix in mathematics

2D matrix in the **main memory**

another 2D matrix in the **main memory**

matcpy-host.cpp

```
// input parameters
unsigned matsize = 4000; // num rows and also num cols
int main(const int argc, const char* argv[]) {
  float* matA = new float[matsize * matsize];
  float* matC = new float[matsize * matsize];
  // kernel processing
  for (register int y = 0; y < matsize; ++y) {
    for (register int x = 0; x < matsize; ++x) {
      register unsigned idx = y * matsize + x;
      matC[idx] = matA[idx];
```

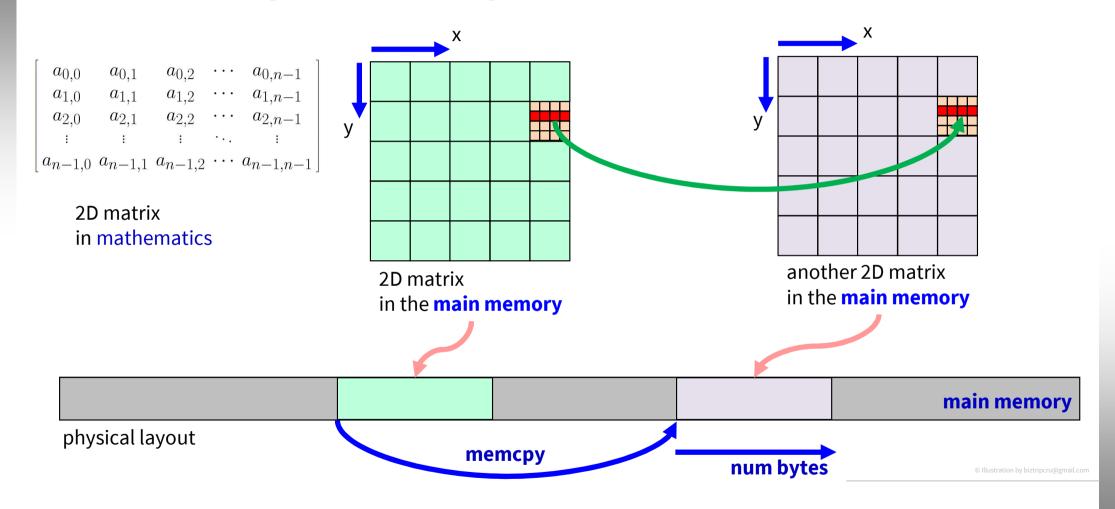
matcpy-host.cpp 실행 결과

434,394 usec for simple 16k x 16k matrix copy (Intel Core i5-3570)

```
linux/cuda-work > ./22a-matcpy-host.exe 16k
elapsed wall-clock time[0] started
elapsed wall-clock time[0] = 434394 usec
matrix size = matsize * matsize = 16384 * 16384
sumA = 134076296.0000000
sumC = 134076296.0000000
diff(sumA, sumC) = 0.0000000
diff(sumA, sumC) / SIZE = 0.000000
matA=[ 0.383000 0.886000 0.777000 ... 0.942000 0.961000 0.764000
        0.991000 0.024000 0.144000 ... 0.318000 0.279000 0.474000
        0.345000 0.967000 0.997000 ... 0.016000 0.090000 0.961000
        0.093000 0.151000 0.150000 ... 0.711000 0.247000 0.182000
        0.395000 0.262000 0.865000 ... 0.435000 0.172000 0.009000
       0.530000 0.136000 0.971000 ... 0.179000 0.510000 0.833000 ]
matC=[ 0.383000 0.886000 0.777000 ... 0.942000 0.961000 0.764000
        0.991000 0.024000 0.144000 ... 0.318000 0.279000 0.474000
        0.345000 0.967000 0.997000 ... 0.016000 0.090000 0.961000
        0.093000 0.151000 0.150000 ... 0.711000 0.247000 0.182000
        0.395000 0.262000 0.865000 ... 0.435000 0.172000 0.009000
        0.530000 0.136000 0.971000 ... 0.179000 0.510000 0.833000 7
linux/cuda-work >
```

CPU version 434,394 usec

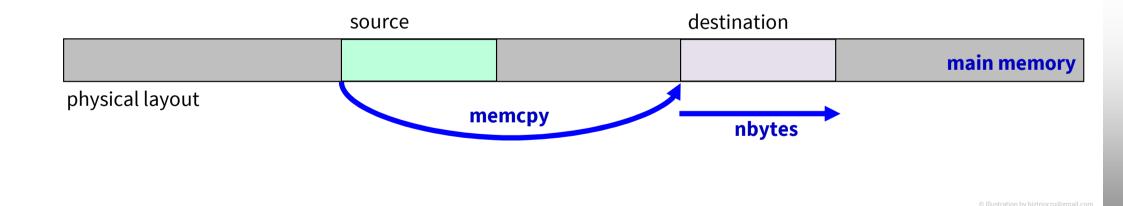
Matrix Copy – memcpy version



memcpy()

void * memcpy(void * destination, const void * source, size_t num);

- main memory 상에서, *source* → destination 으로 *num* byte 를 copy _{복사}
- 겹치는 부분이 있으면, 복사 실패 가능 (memmove() 참고)



matcpy-memcpy.cpp

```
// input parameters
unsigned matsize = 4000; // num rows and also num cols
int main(const int argc, const char* argv[]) {
 // host-side data
  float* matA = new float[matsize * matsize];
  float* matC = new float[matsize * matsize];
  setNormalizedRandomData( matA, matsize * matsize );
  // kernel processing
  ELAPSED_TIME_BEGIN(0);
  memcpy( matC, matA, matsize * matsize * sizeof(float) );
  ELAPSED_TIME_END(0);
```

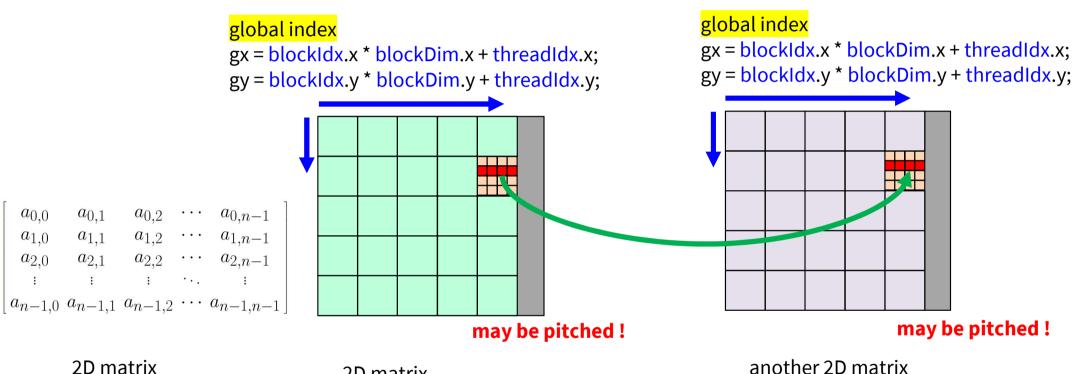
matcpy-memcpy.cpp 실행 결과

449,107 usec for simple 16k x 16k matrix copy (Intel Core i5-3570)

```
linux/cuda-work > ./22b-matcpy-memcpy.exe 16k
elapsed wall-clock time[0] started
elapsed wall-clock time[0] = 449107 usec
matrix size = matsize * matsize = 16384 * 16384
sumA = 134076296.0000000
sumC = 134076296.0000000
diff(sumA, sumC) = 0.000000
diff(sumA, sumC) / SIZE = 0.0000000
matA=[ 0.383000 0.886000 0.777000 ... 0.942000 0.961000 0.764000
        0.991000 0.024000 0.144000 ... 0.318000 0.279000 0.474000
        0.345000 0.967000 0.997000 ... 0.016000 0.090000 0.961000
        0.093000 0.151000 0.150000 ... 0.711000 0.247000 0.182000
        0.395000 0.262000 0.865000 ... 0.435000 0.172000 0.009000
        0.530000 0.136000 0.971000 ... 0.179000 0.510000 0.833000 1
matC=[ 0.383000 0.886000 0.777000 ... 0.942000 0.961000 0.764000
        0.991000 0.024000 0.144000 ... 0.318000 0.279000 0.474000
        0.345000 0.967000 0.997000 ... 0.016000 0.090000 0.961000
        0.093000 0.151000 0.150000 ... 0.711000 0.247000 0.182000
        0.395000 0.262000 0.865000 ... 0.435000 0.172000 0.009000
        0.530000 0.136000 0.971000 ... 0.179000 0.510000 0.833000 ]
linux/cuda-work >
```

CPU version	434,394 usec
memcpy version	449,107 usec

Matrix Copy - CUDA version



2D matrix in mathematics

2D matrix in the **global memory**

in the **global memory**

index 계산

pitched matrix

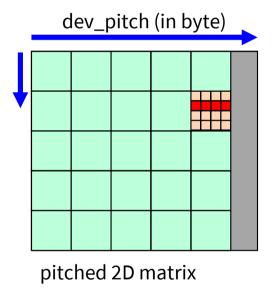
- dev_pitch : cudaMallocPitch(), cudaMalloc3D() → byte 단위
- data : float \rightarrow sizeof(float) = 4 byte

• A[y][x] 의 위치 계산: byte 기준

- offset = y * dev_pitch + x * sizeof(float); // in byte
- *((float*)((char*)A + offset) = 0.0f;

• A[y][x] 의 위치 계산: index 기준

- assert(dev_pitch % sizeof(float) == 0);
- pitch_in_elem = dev_pitch / sizeof(float);
- idx = y * pitch_in_elem + x;
- A[idx] = 0.0f;



matcpy-dev.cu

```
// input parameters
unsigned matsize = 4000; // num rows and also num cols

__global__ void kernelMatCpy( float* C, const float* A, int matsize, size_t pitch_in_elem ) {
    register unsigned gy = blockIdx.y * blockDim.y + threadIdx.y; // CUDA-provided index
    if (gy < matsize) {
        register unsigned gx = blockIdx.x * blockDim.x + threadIdx.x; // CUDA-provided index
        if (gx < matsize) {
            register unsigned idx = gy * pitch_in_elem + gx; // in element
            C[idx] = A[idx];
        }
    }
}</pre>
```

matcpy-dev.cu _{계속}

```
int main(const int argc, const char* argv[]) {
  // allocate device memory
  size_t dev_pitch = 0;
  cudaMallocPitch( (void**)&dev_matA, &dev_pitch, matsize * sizeof(float), matsize );
  cudaMallocPitch((void**)&dev_matC, &dev_pitch, matsize * sizeof(float), matsize);
  // copy to device from host
  cudaMemcpy2D(...);
  // CUDA kernel call
  dim3 dimBlock(32, 32, 1);
  dim3 dimGrid(div_up(matsize, dimBlock.x), div_up(matsize, dimBlock.y), 1);
  assert(dev_pitch % sizeof(float) == 0);
  register unsigned pitch_in_elem = dev_pitch / sizeof(float);
  kernelMatCpy <<< dimGrid, dimBlock >>> ( dev_matC, dev_matA, matsize, pitch_in_elem );
  cudaDeviceSynchronize();
```

matcpy-dev.cu 실행 결과

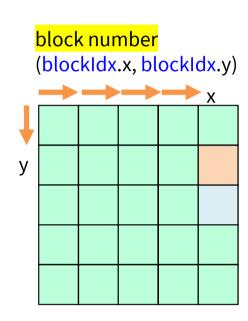
• 6,728 usec for simple 16k x 16k matrix copy (GeForce RTX 2070)

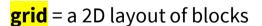
```
linux/cuda-work > ./22c-matcpy-dev.exe 16k
elapsed wall-clock time[] started
dev pitch = 65536 byte, host pitch = 65536 byte
prob size = 16384 * 16384
gridDim = 512 * 512 * 1
blockDim = 32 * 32 * 1
total thr = 16384 * 16384 * 1
elapsed wall-clock time[0] started
elapsed wall-clock time[0] = 6728 usec
elapsed wall-clock time[1] = 857506 usec
matrix size = matsize * matsize = 16384 * 16384
sumA = 134076296.0000000
sumC = 134076296.0000000
diff(sumA. sumC) = 0.000000
diff(sumA, sumC) / SIZE = 0.000000
matA=[ 0.383000 0.886000 0.777000 ... 0.942000 0.961000 0.764000
        0.991000 0.024000 0.144000 ... 0.318000 0.279000 0.474000
        0.345000 0.967000 0.997000 ... 0.016000 0.090000 0.961000
        0.093000 0.151000 0.150000 ... 0.711000 0.247000 0.182000
        0.395000 0.262000 0.865000 ... 0.435000 0.172000 0.009000
       0.530000 0.136000 0.971000 ... 0.179000 0.510000 0.833000 1
matC=[ 0.383000 0.886000 0.777000 ... 0.942000 0.961000 0.764000
        0.991000 0.024000 0.144000 ... 0.318000 0.279000 0.474000
```

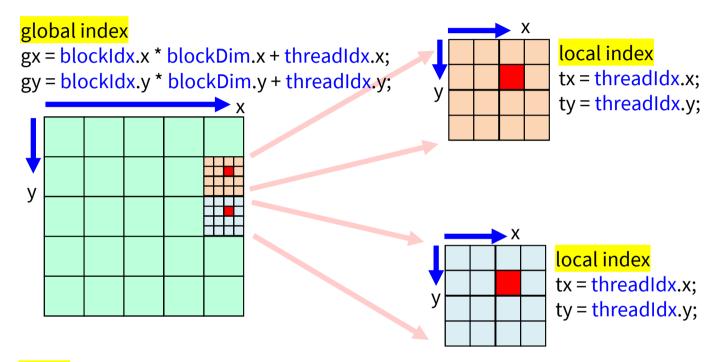
CPU version	434,394 usec
memcpy version	449,107 usec
CUDA naïve copy	6,728 usec

2216

2D Tiled Layout for Matrix Operations





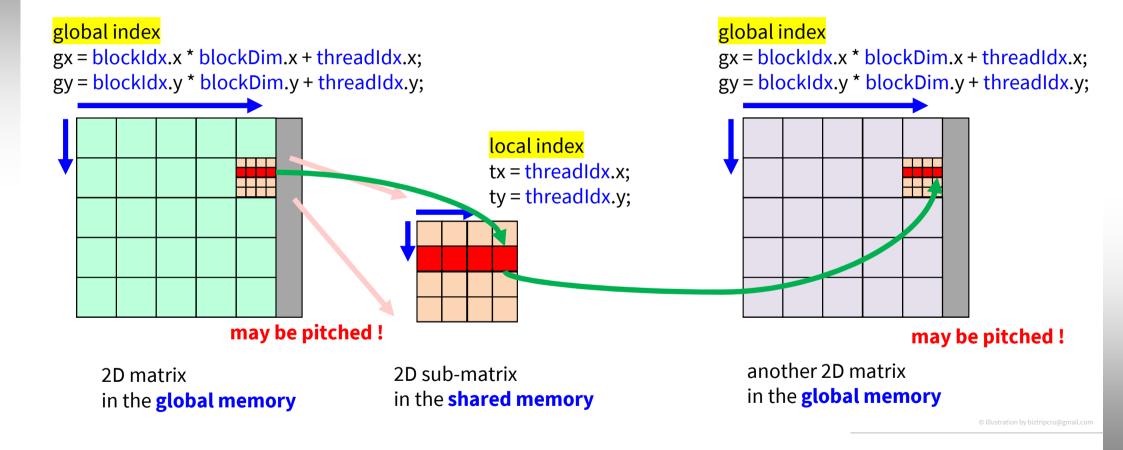


block = a 2D layout of threads

thread block in **shared memory**

© Illustration by biztripcru@gmail.co

Matrix Copy – Tiled Approach



Tile Size

- for simplicity, we assume square matrices
- tile : 정사각형 모양을 선호

```
<u>__shared</u>__ float s_mat[32][32];
```

- $32 \times 32 = 1,024$
- 1,024 = maximum number of threads in a thread block
- 32 = a single warp
- old code 에서는 16x16 또는 32x16 사용
 - 이유: maximum number of threads in a thread block was 512
 - 16 x 16 = 256
 - 32 x 16 = 512 → 1개의 thread 에서 2개씩 처리 → 32x32 로 작동

matcpy-shared.cu

```
// CUDA kernel function
  _global___ void kernelMatCpy( float* C, const float* A, unsigned matsize, <mark>size_t pitch_in_elem</mark> ) {
   _shared__ float s_mat[32][32];
  register unsigned gy = blockIdx.y * blockDim.y + threadIdx.y; // CUDA-provided index
  if (gy < matsize) {</pre>
    register unsigned gx = blockIdx.x * blockDim.x + threadIdx.x; // CUDA-provided index
    if (gx < matsize) {</pre>
      register unsigned idx = gy * pitch_in_elem + gx; // in element
      s_mat[threadIdx.y][threadIdx.x] = A[idx];
      __syncthreads();
      C[idx] = s_mat[threadIdx.y][threadIdx.x];
```

matcpy-shared.cu 실행 결과

• 7,364 usec for simple 16k x 16k matrix copy, shared memory (GeForce RTX 2070)

```
linux/cuda-work : ./22d-matcpy-shared.exe 16k
                                                                      CPU version
                                                                                                 434,394 usec
elapsed wall-clock time[1] started
                                                                      memcpy version
                                                                                                 449.107 usec
dev pitch = 65536 byte, host pitch = 65536 byte
                                                                      CUDA naïve copy
                                                                                                   6,728 usec
prob size = 16384 * 16384
gridDim = 512 * 512 * 1
                                                                      CUDA shared mem copy 7,364 usec
blockDim = 32 * 32 * 1
total thr = 16384 * 16384 * 1
elapsed wall-clock time[0] started
elapsed wall-clock time[0] = 7364 usec
                                                         slower than global memory version!
elapsed wall-clock time[1] = 866073 usec
matrix size = matsize * matsize = 16384 * 16384
sumA = 134076296.0000000
sumC = 134076296.0000000
diff(sumA, sumC) = 0.000000
diff(sumA, sumC) / SIZE = 0.000000
matA=[ 0.383000 0.886000 0.777000 ... 0.942000 0.961000 0.764000
       0.991000 0.024000 0.144000 ... 0.318000 0.279000 0.474000
       0.345000 0.967000 0.997000 ... 0.016000 0.090000 0.961000
       0.093000 0.151000 0.150000 ... 0.711000 0.247000 0.182000
       0.395000 0.262000 0.865000 ... 0.435000 0.172000 0.009000
       0.530000 0.136000 0.971000 ... 0.179000 0.510000 0.833000 ]
matC=[ 0.383000 0.886000 0.777000 ... 0.942000 0.961000 0.764000
                                                                                                           2221
       0.991000 0.024000 0.144000 ... 0.318000 0.279000 0.474000
```

matcpy-cudacpy.cu

```
int main(const int argc, const char* argv[]) {
 // CUDA kernel call
 ELAPSED_TIME_BEGIN(0);
 cudaMemcpy2D( dev_matC, dev_pitch, dev_matA, dev_pitch,
                                        matsize * sizeof(float), matsize, cudaMemcpyDeviceToDevice);
 cudaDeviceSynchronize();
 ELAPSED_TIME_END(0);
 CUDA_CHECK_ERROR();
```

matcpy-cudacpy.cu 실행 결과

• 5,656 usec for simple 16k x 16k matrix copy, cudaMemcpy2D() (GeForce RTX 2070)

```
linux/cuda-work > ./22e-matcpy-cudacpy.exe 16k
elapsed wall-clock time[1] started
dev pitch = 65536 byte, host pitch = 65536 byte
elapsed wall-clock time[0] started
elapsed wall-clock time[0] = 5656 usec
elapsed wall-clock time[1] = 854080 usec
matrix size = matsize * matsize = 16384 * 16384
sumA = 134076296.0000000
sumC = 134076296.000000
diff(sumA, sumC) = 0.000000
diff(sumA, sumC) / SIZE = 0.000000
matA=[ 0.383000 0.886000 0.777000 ... 0.942000 0.961000 0.764000
        0.991000 0.024000 0.144000 ... 0.318000 0.279000 0.474000
        0.345000 0.967000 0.997000 ... 0.016000 0.090000 0.961000
        0.093000 0.151000 0.150000 ... 0.711000 0.247000 0.182000
        0.395000 0.262000 0.865000 ... 0.435000 0.172000 0.009000
        0.530000 0.136000 0.971000 ... 0.179000 0.510000 0.833000 ]
matC=[ 0.383000 0.886000 0.777000 ... 0.942000 0.961000 0.764000
        0.991000 0.024000 0.144000 ... 0.318000 0.279000 0.474000
        0.345000 0.967000 0.997000 ... 0.016000 0.090000 0.961000
        0.093000 0.151000 0.150000 ... 0.711000 0.247000 0.182000
        0.395000 0.262000 0.865000 ... 0.435000 0.172000 0.009000
        0.530000 0.136000 0.971000 ... 0.179000 0.510000 0.833000 ]
linux/cuda-work >
```

CPU version	434,394 usec
memcpy version	449,107 usec
CUDA naïve copy	6,728 usec
CUDA shared mem copy	7,364 usec
CUDA memcpy2D	5,656 usec

내용contents

• matrix copy - theoretical limit

CPU version	434,394 usec
memcpy version	449,107 usec
CUDA naïve version – global memory	6,728 usec
CUDA shared memory version	7,364 usec
CUDA memcpy2D	5,656 usec

Matrix Copy

행렬 복사

폰트 끝단 일치 → 큰 교자 타고 혼례 치른 날 정**참판 양반댁 규수 큰 교자 타고 혼례 치른 날** 정 참판 양반댁 규수 큰 교자 타고 혼례 치른 날 본고딕 Noto Sans KR

© 2021-2022. biztripcru@gmail.com. All rights reserved. 모든 저작권은 biztripcru@gmail.com 에게 있습니다.



Mathematical Notations $O(n \log n)$ **Source Serif Pro**

Source Sans Pro